The Invisible Costs of AI: Global Data Labour and the Limits of Ethical Governance

Author- Anurita Srivastava
Affiliation- Georgetown University, McCourt School of Public Policy
Program- Master of Public Policy (MPP)

Introduction

Artificial intelligence (AI) is widely promoted as a public good, an efficiency-enhancing tool that can modernize governance, expand access to services, and drive economic growth. Governments describe AI as essential infrastructure for the future, while technology firms frame their products as neutral, scalable solutions that benefit humanity as a whole. This optimism has helped position AI as both inevitable and desirable, often sidelining questions about how these systems are actually built.

That framing breaks down when we look closely at the labour that makes AI work. In Nairobi, a 28-year-old data labeller, referred to here as Fasica, spends up to nine hours a day reviewing graphic videos of violence and abuse so that AI systems can learn what content to block or flag. For this work, she earns roughly $1.50 an hour (Ajuzieogu 2025). Like thousands of workers across Kenya, India, Ghana, and the Philippines, Fasica performs tasks essential to the functioning of “safe” AI systems, yet she remains largely invisible in policy debates about AI governance.

This essay argues that the global AI labour economy represents a systemic ethical failure in public policy. Using Fasica’s experience as a throughline, alongside interviews with members of the Data Labelers Association (DLA) and existing empirical research, the essay shows how AI governance frameworks externalize psychological harm and economic risk onto workers in the Global South while concentrating profits and decision-making power in the Global North. Rather than treating these harms as unfortunate side effects, I argue that they are predictable outcomes of policy choices embedded in current governance regimes. A more just approach to AI governance requires recognizing data infrastructure workers as core stakeholders in technological development and redesigning policies to distribute both benefits and burdens more fairly across borders.

The Hidden Labour That Trains AI

Despite popular narratives of automation, AI systems depend heavily on human labour. Since algorithms do not learn on their own, they are trained through millions of human judgments. Data labelers classify images, annotate text, transcribe audio, and review disturbing content to ensure AI systems function reliably. This work is typically outsourced through layers of subcontractors to lower-income countries, allowing major technology firms to distance themselves from labour conditions while retaining control over profits and intellectual property (Gray and Suri 2019).

Fasica's experience reflects the human cost of this structure. Her job involves repeated exposure to extreme violence, hate speech, and sexual abuse material. According to interviews conducted with DLA members, workers are expected to meet strict productivity targets, often reviewing hundreds of pieces of content per day with little time to recover emotionally. Training on mental health risks is minimal, and long-term psychological support is rare or nonexistent.[1] Several workers described

1 **Data Labelers Association**, interview with the Kenyan association founded by a group of data labelers advocating for fair treatment, better conditions, and recognition worldwide. Unpublished interview conducted by the author, 2025.

symptoms consistent with trauma, including insomnia, anxiety, emotional numbness, and intrusive thoughts.[2]

These harms are not accidental. They are built into a system designed to keep labour cheap, replaceable, and geographically distant from end users. Wages remain low, contracts are short-term, and termination can be sudden. Workers reported constant fear of replacement, limited access to paid leave, and few avenues for legal recourse. The contrast between the multi-billion-dollar valuations of AI firms and the insecurity faced by workers like Fasica highlights a stark ethical imbalance: those who bear the most significant risks receive the smallest share of the benefits.

Efforts to organize or raise concerns often trigger retaliation. DLA members reported informal blocklisting across contractors, effectively shutting workers out of future employment. Importantly, these workers do not oppose AI itself. Their demands are modest and concrete: fair pay, mental health protections, job security, and a voice in decisions that shape their working conditions. Their exclusion from governance processes exposes the gap between AI's universalist rhetoric and the unequal realities of its production.

Policy Assumptions That Normalize Harm

---

2 For empirical evidence on psychological harm, trauma exposure, and long-term mental health impacts among content moderators and data workers, see Sarah T. Roberts, *Behind the Screen: Content Moderation in the Shadows of Social Media* (New Haven, CT: Yale University Press, 2019); and Mary L. Gray and Siddharth Suri, *Ghost Work: How to Stop Silicon Valley from Building a New Global Underclass* (Boston: Houghton Mifflin Harcourt, 2019).

Current AI governance frameworks rest on several implicit assumptions that help explain why data labour remains politically invisible. First, there is a strong emphasis on aggregate benefits. Policymakers often assume that if AI increases overall efficiency or welfare, its development is morally justified. This logic mirrors a utilitarian approach that prioritizes outcomes over the distribution of harms (Mill 1861). As a result, regulation tends to focus on downstream risks, such as algorithmic bias or consumer safety, while ignoring upstream labour conditions.

Second, policymakers frequently rely on the idea that market contracts signal valid consent. If workers agree to the job, the reasoning goes, then the arrangement is ethically acceptable. However, consent formed under conditions of deep structural inequality is morally thin.[3] For Fasica, declining this work does not mean choosing a better job; it often means choosing unemployment. Treating such constrained choices as genuine consent allows governments to frame exploitative labour practices as private-market outcomes rather than public-policy concerns (Wolff 2023).

Third, responsibility is fragmented across global supply chains. Major AI firms outsource labour to contractors, who in turn rely on subcontractors, creating plausible deniability at every level. This reflects as a failure to address structural injustice- harm persists precisely because no single actor is held fully responsible (Young 2011). Together, these assumptions produce a governance architecture in which harm is foreseeable, routine, and yet institutionally ignored.[4]

3 For critiques of consent under conditions of structural inequality and constrained choice, see Jonathan Wolff, *An Introduction to Political Philosophy*, 4th ed. (Oxford: Oxford University Press, 2023), chap. 7.

4 This pattern aligns with Iris Marion Young’s account of structural injustice, in which harm persists not because of isolated wrongdoing but because responsibility is dispersed across institutions and actors in ways that undermine accountability. See Iris Marion Young, *Responsibility for Justice* (Oxford: Oxford University Press, 2011).

Distributive and Cosmopolitan Justice

From a distributive justice perspective, the global AI labour economy is challenging to defend. Rawls's difference principle allows inequality only when it benefits the least advantaged (Rawls 1971). In the AI sector, the opposite pattern holds. Psychological harm and economic insecurity are concentrated among already-marginalized workers, while gains accrue to firms, consumers, and governments in wealthy countries.

A cosmopolitan approach strengthens this critique by rejecting the idea that moral obligations stop at national borders. AI systems are global by design, as they rely on transnational labour, generate cross-border profits, and shape information environments worldwide. If the benefits of AI are global, then so too are the responsibilities (Pogge 2006). Workers like Fasica are not morally peripheral simply because they live outside the jurisdictions of the AI firms' headquarters.

Policy frameworks that treat data labeling as low-skilled, interchangeable work further entrench injustice. By excluding this labour from classifications that trigger stronger occupational protections, governments effectively render psychological trauma an acceptable cost of innovation. This is not a neutral outcome of technology but a direct result of regulatory choices that prioritize speed and competitiveness over fairness.

Legitimacy, Consent, and the Social Contract

Social contract theory raises further doubts about the legitimacy of current AI governance.[5] Political and economic arrangements are legitimate only when they can be reasonably accepted as fair terms of cooperation (Rawls 1971). While data workers technically agree to their contracts, extreme power asymmetries undermine meaningful reciprocity.

Fasica has no say in how the data she labels are used, how profits are distributed, or how risks are managed. Benefits flow upward; harms flow downward. Workers are treated as inputs rather than participants in a shared project of technological development. Under these conditions, compliance does not equal legitimacy. An arrangement that extracts labour without fair participation or shared benefit fails basic standards of moral justification (Scanlon 1998; Sen 2009).

Counterargument

A common defense of the AI labour economy is that data labeling provides income where few alternatives exist. From this perspective, even low-paid or psychologically harmful jobs are preferable to unemployment, and stronger regulation might eliminate these opportunities.

This argument is ethically weak.[6] First, it treats background injustice as a justification for further harm. The fact that Fasica has limited options does not make exposing her to trauma morally acceptable. Second, it sets a dangerously low bar- exploitation becomes permissible whenever workers are poor enough. Finally, it ignores the active role firms play in designing systems to minimize labour costs while maximizing profit.

5 For legitimacy understood as reasonable acceptance and fair terms of cooperation, see T. M. Scanlon, *What We Owe to Each Other* (Cambridge, MA: Harvard University Press, 1998).
6 For the argument that justice evaluates fairness rather than comparisons to the worst available alternative, see Martha C. Nussbaum, *Frontiers of Justice* (Cambridge, MA: Harvard University Press, 2006).

Justice does not ask whether an arrangement is better than the worst alternative; it asks whether it is fair. An economy that relies on trauma because it is cheap is not redeemed by the absence of better options.

Precedents for Ethical Regulation

Regulating AI labour is not unprecedented. Other policy domains already recognize ethical responsibility for transnational harm. Conflict minerals regulation, such as Section 1502 of the U.S. Dodd-Frank Act and the EU Conflict Minerals Regulation, requires firms to conduct due diligence on global supply chains when production is linked to violence and exploitation (Source Intelligence 2022). These frameworks explicitly reject the idea that distance or subcontracting absolves responsibility.[7]

Environmental regulation offers a second parallel through the polluter-pays principle[8], which holds firms responsible for the social costs of production (OECD 1972). In the AI context, psychological trauma functions as an unpriced externality. Recognizing it as a policy-relevant cost would align AI governance with long-standing regulatory norms rather than representing a radical departure.

Towards Just AI Governance

7 See also *OECD Guidelines for Multinational Enterprises* (Paris: OECD Publishing, 2011; updated 2023), which extend corporate due diligence and responsibility across global supply chains.

8 On the polluter-pays principle (PPP) as a cornerstone of regulatory ethics, see OECD, *Recommendation of the Council on Guiding Principles Concerning International Economic Aspects of Environmental Policies* (Paris: OECD, 1972).

A more just approach to AI governance starts by recognizing data workers as essential contributors rather than disposable inputs. Content moderation and data labeling should be classified as high-risk labour, triggering mandatory mental health protections, workload limits, hazard pay, and long-term healthcare coverage. Psychological harm should be treated as a legitimate workplace injury, not an inevitable byproduct of innovation.

Economic mechanisms can further rebalance benefits and burdens. Data labour royalties would tie worker compensation to the commercial success of AI systems, acknowledging their ongoing contributions. A psychological harm levy could require firms to internalize the costs of trauma exposure, discouraging business models that rely on constant exposure to extreme content.

At the global level, AI governance institutions should include formal representation from data-labour communities in the Global South. A co-governance body with agenda-setting authority over labour standards would move beyond symbolic consultation and challenge the assumption that AI policy is the exclusive domain of Northern states and corporations.

Conclusion

Fasica’s labour helps make AI systems safer and more usable for millions of people who will never know her name. However, current governance frameworks treat her suffering as irrelevant to the moral evaluation of AI. This disconnect reveals a more profound crisis in how technological progress is governed.

Ethical AI cannot be achieved by focusing solely on more innovative algorithms or better outputs. It must begin with justice for the people whose labour sustains these systems. When workers like Fasica are recognized as stakeholders with rights, protections, and a voice, innovation becomes not only more ethical but more legitimate. Only through this lens of justice can we create a future where technology thrives alongside ethical considerations, ultimately contributing to a sustainable and just society that values all individuals equally. Without this shift, AI governance will continue to reproduce global inequalities while claiming to serve humanity as a whole.

Word count- 1753

References

- Ajuzieogu, Uche. 2025. "The Hidden Price of Every ChatGPT Query: How Streaming AI Built an Empire on $1.50 an Hour." *Aylgorith*, October 25, 2025. https://aylgorith.com/streaming-ai-economics/.
- Beitz, Charles R. 1979. *Political Theory and International Relations*. Princeton, NJ: Princeton University Press. https://archive.org/details/politicaltheory
- Berg, Janine, Marianne Furrer, Ellie Harmon, Uma Rani, and M. Six Silberman. 2018. *Digital Labour Platforms and the Future of Work*. Geneva: International Labour Organization.
- Birhane, Abeba. 2020. "Algorithmic Colonization of Africa." *SCRIPTed* 17 (2): 389–409. https://doi.org/10.2966.
- Cohen, G. A. 2006. "Where the Action Is: On the Site of Distributive Justice." *Philosophy & Public Affairs* 26 (1): 3–30. https://onlinelibrary.wiley.com/doi/10.1111

- Corporate Sustainability Due Diligence Directive (CSDDD). 2024. Directive (EU) 2024/1760. EUR-Lex. https://eur-lex.europa.eu/eli.
- Data Labelers Association. 2025. *Interview with Data Labeling and Content Moderation Workers*. Unpublished interview conducted by the author.
- Gray, Mary L., and Siddharth Suri. 2019. *Ghost Work: How to Stop Silicon Valley from Building a New Global Underclass*. Boston: Houghton Mifflin Harcourt. *https://labordoc.ilo.org/*
- Kant, Immanuel. 1785. *Groundwork of the Metaphysics of Morals*. Various editions.
- Mill, John Stuart. 1861. *Utilitarianism*. London: Parker, Son, and Bourn.
- Nussbaum, Martha C. 2006. *Frontiers of Justice: Disability, Nationality, Species Membership*. Cambridge, MA: Harvard University Press.
- OECD. 1972. *Recommendation of the Council on Guiding Principles Concerning International Economic Aspects of Environmental Policies*. Paris: OECD. https://legalinstruments.oecd.org/en/instruments/OECD-LEGAL-0102.
- OECD. 2011. *OECD Guidelines for Multinational Enterprises*. Paris: OECD Publishing. Updated 2023. https://www.oecd.org/corporate/mne/.
- Pogge, Thomas. 2006. "World Poverty and Human Rights." *Ethics & International Affairs* 19 (1): 1–7. *https://www.researchgate.net/publication/World_Poverty_and_Human_Rights*
- Rawls, John. 1971. *A Theory of Justice*. Cambridge, MA: Harvard University Press.
- Roberts, Sarah T. 2019. *Behind the Screen: Content Moderation in the Shadows of Social Media*. New Haven, CT: Yale University Press.
- Rousseau, Jean-Jacques. 1762. *The Social Contract*. Various editions.

- Scanlon, T. M. 1998. *What We Owe to Each Other*. Cambridge, MA: Harvard University Press.
- Sen, Amartya. 2009. *The Idea of Justice*. Cambridge, MA: Harvard University Press.
- Source Intelligence. 2022. "EU Conflict Minerals Regulation vs. U.S. Dodd-Frank Act Section 1502." August 8, 2022. https://blog.sourceintelligence.com/blog/eu-conflict-minerals-vs-us-dodd-frank-act-section-1502.
- Wolff, Jonathan. 2023. *An Introduction to Political Philosophy*. 4th ed. Oxford: Oxford University Press.
- Young, Iris Marion. 2011. *Responsibility for Justice*. Oxford: Oxford University Press. https://academic.oup.com/book/4381